

A FOOL'S
DREAM

HUMANITY UNDER TENSION

AI FOR
REPLACEMENT

MAN
NOT
MACHINE

OBSOLETE

The Tale of Two AIs

Artificial and Augmented

AI FOR AUGMENTATION

Basil C. Puglisi, MPA

T H E

Tale of Two AIs

Artificial and Augmented

Basil C. Puglisi, MPA

A Human-AI Collaboration · basilpuglisi.com

May 2026



This is the age of replacement, and this is the age of augmentation. This is the season when AI removes people from the workflow, and this is the season when AI raises the capability of the people who remain. This is the quarter when executives fund infrastructure through headcount reduction, and this is the quarter when other executives use the same systems to cut friction, preserve judgment, and expand capacity. This is the moment of cognitive offloading, and this is the moment of cognitive development.

Two cultures are being built from the same code. One trains workers for dependency and calls it efficiency. The other trains workers for judgment and calls it growth. Both cultures claim the future, but only one keeps the human being as a developing participant in the system. There is talk of a Third Culture, but it never coalesces, because it never has. The prohibition dream is a fool's dream: STOPAI, close the box, return the genie to the bottle.

Consider two of those decisions, made in the same year, deploying the same class of technology, choosing opposite architectures.

On March 31, 2026, Oracle began a major restructuring reported to involve thousands of employees, while building toward an analyst-estimated \$156 billion in multi-year AI infrastructure commitments.¹ The accompanying \$2.1 billion restructuring charge and the redirection of \$8 to \$10 billion in annual cash flow from payroll toward data centers and GPU clusters established the strategic logic. Oracle's own March 2026 10-Q filing, submitted to the SEC, confirmed the direction in primary language: "AI models for generating computer code have become so efficient that we have been restructuring our product development teams into smaller, more agile and productive groups. This new AI Code Generation technology is enabling us to build more software in less time with fewer people."¹

The secondary reporting around the restructuring carried the executive framing. At a Stargate investor briefing on April 2, 2026, Larry Ellison was reportedly quoted as saying: "We are choosing the chips." Safra Catz, then Oracle's Executive Vice Chair, was quoted in the same reporting period as framing the announcement as "a generational reallocation of capital from people-intensive consulting and legacy support toward GPU-intensive AI infrastructure." The announcement functions in this paper as the clearest named instance of a pattern that had been building across the technology sector for the prior eighteen months, and it gave the pattern a canonical statement of intent.

The same frontier model generation that powered Oracle's investment decision also powered Cynergy Bank's deployment decision. Cynergy embedded AI into customer-facing operations not to reduce headcount but to extend human

relationship management at scale. Following the app's launch in late October 2024, the deployment produced an 18 percent reduction in customer complaints and a 15 percent reduction in transactional call volume against comparable channel volumes, recorded over twelve months of operation.² The human contact layer expanded rather than contracted. The same technology window, the same class of capability, and the opposite architectural choice: Oracle chose the chips; Cynergy chose the people who use them better.

This paper doesn't argue that Oracle was wrong and Cynergy was right in the sense of strategic miscalculation. The argument is structural. Two architectures are operationally available using the same frontier models, and the choice between them is consequential in ways that compound across time. The first architecture treats AI as a replacement for human labor, and the second treats AI as infrastructure for human capability. The first produces what this paper calls Culture 1; the second produces Culture 2. A third option, prohibition of the technology, is not operationally available for foundational reasons examined in the final section. The paper documents both operational architectures, explains why the third is unavailable, and examines the governance infrastructure that determines which choice an institution makes.

PLATE I — THE CHOOSING OF THE CHIPS.



THE REPLACEMENT ARCHITECTURE

The state-scale exemplar

The People's Republic of China has chosen the replacement-plus-control architecture most openly and built the regulatory, infrastructural, and ideological scaffolding to sustain it at maximum institutional scale. This is not a claim about China as an adversary. It is a factual observation: China makes the Culture 1 architecture visible as a sustained state-level commitment. What other states and corporations have chosen partially, ambiguously, or under cover, China has chosen in policy text, binding regulation, and named enforcement.

The Cyberspace Administration of China governs AI deployment through a three-pillar framework that has accumulated binding regulatory authority across the period from 2022 through 2026. Algorithm regulation, deep synthesis provisions, and the Interim Measures for the Management of Generative AI Services (effective August 15, 2023, the first administrative regulation on generative AI services globally) together require that generative AI services with public-opinion or social-mobilization capabilities complete a Dual Filing process: mandatory pre-launch security assessment with provincial CAC offices plus algorithmic filing under separate provisions. Training data must align with state-defined value criteria. As of March 2025, 346 generative AI services had completed mandatory LLM filing; the remainder faced enforcement.

Mandatory AIGC content labeling took effect September 1, 2025, requiring explicit user-visible labels and implicit metadata-embedded labels across text, images, audio, video, and virtual scenes. Three national security standards governing data annotation, pre-training and fine-tuning data, and baseline security requirements followed on November 1, 2025. The Qinglang 2026 enforcement campaign, launched by CAC in April 2026, introduced mandatory labeling and quarterly algorithmic audits for high-impact services. The Interim Measures for Anthropomorphic AI Interactive Services, published April 10, 2026 and effective July 15, 2026, extend regulation to AI chatbots, companions, and customer service agents. Draft rules governing digital virtual human services, published April 3, 2026, address biometric deepfakes and informed consent.

The deployment pattern within this regulatory envelope shows state-directed AI distributed across MedTech (DeepSeek and HuatuoGPT deployed in public hospitals for diagnostics, imaging, intelligent triage, and chronic disease management), retail and e-commerce (algorithmic content and payment

architectures with state-aligned recommendation requirements), and manufacturing automation (AI-driven robotics absorbing inspection, warehousing, and perimeter-patrol roles at Foxconn and Shenzhen-adjacent scale). Some of these deployments carry augmentation characteristics at the application layer; the claim is not that every AI deployment in China is replacement-pattern, but that the governing architecture is replacement-plus-control at state scale.

The information architecture is integrated by design. The Great Firewall as foundational layer, algorithmic recommendation regulations requiring state-aligned outputs, deep synthesis rules, and mandatory content labeling together produce a stack in which value alignment with state-defined criteria is a binding deployment requirement rather than an advisory standard. WeChat functions as a state-monitored daily-life platform across the full population. Digital identity verification centralizes user data flow to state repositories. The Carnegie Endowment's July 2025 analysis described this configuration as a Crossroads Era post-DeepSeek,³ with the party asserting control over AI outputs, user data, and safety frameworks simultaneously. The analysis named the AI Plus plan, adopted at the 2024 Central Economic Work Conference, as the integration of AI into the real economy through state direction, with export-economy applications in manufacturing robotics and BYD-scale automation as the competitive mechanism.

The point is not that China regulates AI more aggressively than other states. The point is that China binds capability, data, labeling, and state alignment into a single deployment system. What China has chosen openly and at state scale, individual firms and other governments have chosen partially and less explicitly. The corporate-scale evidence in the following sections documents the Culture 1 architecture operating under Western institutional framing.

The named architecture

Replacement deployment treats AI as substitution for human labor rather than extension of human capability. The deployment pattern produces consequences in a predictable sequence: headcount reductions funded by AI investment, training-pipeline collapse as entry-level positions disappear, cognitive offloading as the work that builds expertise is absorbed by AI before practitioners can develop the expertise needed to govern it, hollowed career ladders as the mid-level roles that would have been built from entry-level experience fail to form, and Universal Basic Income as the remaining social contract in the policy conversation once the workforce pipeline has collapsed.

Three bodies of evidence document this architecture as currently operational and producing its predicted consequences at scale.

The headcount evidence

The year 2025 produced 154,445 technology-sector layoff announcements in the United States, up 15 percent from the prior year, with the Challenger, Gray and Christmas tracking report attributing the acceleration to technology firms "pivoting to both developing and implementing AI much more quickly than any other industry."⁴ Through April 2026, Challenger reported 49,135 AI-attributed cuts, "nearly as much as the total for all AI-related layoffs the firm reported in 2025." The Skillsyncer tracker through May 12, 2026 logged 179 layoff events and 113,863 workers affected, at an average of approximately 856 per day.⁵

The named-executive evidence at the corporate tier is specific. Oracle's March 31 announcement, with the Ellison and Catz framings already on record, is the most clearly attributed. Block's Jack Dorsey committed to a 40 percent workforce reduction with an explicit framing of AI capability as the enabling condition. Snap cited AI advancements directly in its restructuring rationale. Atlassian and PayPal joined the same documentation period with comparable framing. Meta's 8,000 performance-management exits, Microsoft's 9,000 buyout-driven departures, and Amazon's 16,000 corporate-role reductions all occurred in 2025 and 2026, though the attribution discipline in those cases is more complex and addressed in the following section.

Andy Challenger, Chief Revenue Officer at Challenger, Gray and Christmas, offered the framing that resolves the attribution question at the systemic level without requiring individual case certainty: "Regardless of whether individual jobs are being replaced by AI, the money for those roles is." The capital is redirecting from payroll to AI infrastructure, and the workforce consequences follow the capital, whether or not any given organization names AI as the cause of any given departure.

The attribution discipline

Not all AI-attributed layoffs are AI-caused. The paper applies a three-tier discipline before treating any layoff as Culture 1 evidence.

Tier A covers named-executive primary attribution: a senior executive explicitly names AI capability as the enabling condition or causal factor in the workforce reduction. Oracle, Block, and Snap meet this standard. These cases are the primary table.

Tier B covers AI-adjacent restructuring: the workforce change occurs in close temporal proximity to AI investment announcements and is plausibly AI-enabled but independently explained or disclaimed by leadership. Microsoft's restructuring, in which Satya Nadella simultaneously described a mission transition from "software factory to intelligence engine" and denied AI as the

direct cause of departures, is the canonical Tier B case. The denial and the mission statement sit in tension, and the paper does not resolve them; it records both the restructuring and the denial as part of the evidence base.

Tier C covers independently denied attribution: organizations that publicly, specifically, and without qualification deny AI as a factor in workforce reductions. Cases in this tier are excluded from the primary evidence base; the denial is preserved as dissent.

Sam Altman's February 2026 acknowledgment of "AI-washing" as a named phenomenon, in which companies attribute workforce reductions to AI capability as a strategic positioning move rather than a factual account, establishes that the attribution discipline is not the paper's invention. The phenomenon has been named by the chief executive of the most prominent organization deploying the technology. Ben May of Oxford Economics framed the same dynamic from the analyst side: "We suspect some firms are trying to dress up layoffs as a good news story rather than a bad one, for example, by pointing to technological change instead of past overhiring."

The attribution discipline narrows the strictly AI-caused layoff base. The remaining narrowed set is still substantial, and the pattern is still operational at named-executive primary attribution. The cognitive and pipeline evidence follows without requiring the disputed cases.

The cognitive dimension

The cognitive consequences of replacement deployment are the most contested element of the Culture 1 evidence base, and the paper applies strict epistemic discipline. The causal claim, that AI-dependent workflows produce measurable reductions in human cognitive capability, is mechanistically plausible but not yet longitudinally validated at population scale.

The MIT Media Lab study published in arXiv (Kosmyna et al., arXiv 2506.08872)⁶ used EEG measurement to document connectivity reductions of up to 55 percent in participants who completed writing tasks with large language model assistance compared to participants who completed the same tasks without AI or with internet-only access. The finding is specific, instrumented, and published as an arXiv preprint pending peer review. It is also subject to a familiarization-effect critique: participants accustomed to AI-assisted writing may exhibit different connectivity patterns from participants encountering AI assistance for the first time, and the study's baseline conditions do not fully control for this possibility. Gerlich (2025, Societies)⁷ used a correlational survey design with 666 United Kingdom participants to document associations between AI use, cognitive offloading, and self-reported reductions in critical thinking ($r =$

0.72 and $r = -0.75$ respectively). The study establishes the association; it does not establish the causal direction.

The paper has separately published a full methodological audit of the cognitive decline narrative, available at basilpuglisi.com. That audit found that the peer-reviewed evidence base does not yet support the causal claim at the standards adjacent fields require for cognitive intervention research. The same audit found that the same evidence base also does not support the opposite claim, that unstructured AI use is cognitively neutral or beneficial. Both directional claims run ahead of available evidence.

What is verifiable without the longitudinal data is the underlying pattern. The Klarna reversal operates as a documented lagging indicator. Klarna's chief executive announced in February 2024 that AI could perform the work of 700 customer service agents; the company expanded its AI-first operational model through 2024. By late 2025 and into 2026, the company began reversing course. Quality deterioration in AI-handled customer interactions became visible approximately 12 to 18 months after deployment, and Klarna moved toward rehiring and hybrid human-AI models. The reversal does not prove cognitive decline at the individual worker level. It shows that replacement architecture can produce quality consequences at the organizational level that become measurable at a lag, and that the organizations which most aggressively implemented replacement deployment are among the first to discover that the quality consequences exceed the efficiency gains.

The training-pipeline collapse

The most consequential element of the Culture 1 evidence base is the cohort-level employment data for early-career workers in AI-exposed occupations.

The Stanford Digital Economy Lab published "Canaries in the Coal Mine: Six Facts about the Recent Employment Effects of Artificial Intelligence" in August 2025⁸, drawing on ADP payroll data covering millions of workers. The study found a 16 percent relative employment decline for workers aged 22 to 25 in the most AI-exposed occupations after controlling for firm-level shocks, with software engineering entry-level positions declining approximately 20 percent and customer service positions declining approximately 11 percent. A follow-up note published in March 2026 qualified the timing attribution: under the broadest set of controls including firm-time fixed effects, the employment decline became statistically significant only in 2024, with the authors noting that some portion of the earlier decline was attributable to the end of zero-interest-rate policy and the broader technology-sector contraction. The directional finding holds across all specifications; the attribution of specific timing entirely to AI does not.

Supplementary data from LinkedIn through 2026 showed AI-related job postings up 340 percent since 2024 and traditional software engineering roles down 15 percent over the same period. The gap between AI-related demand and traditional-role decline represents the compositional shift the Canaries study measures at the payroll level.

The academic dissent register is preserved here as required. Chandar's follow-up analytical work, the Gimbel analysis, the Eckhardt/Goldschlag findings, and the Humlum/Vestergaard study of Danish administrative payroll data, which found at most small changes in earnings and hours across occupations with confidence intervals ruling out effects larger than 1 to 2 percent, all point to a more contested empirical picture than the Stanford headline finding implies. The paper does not resolve this contest; it documents the directional finding and the challenge to it, and carries both forward as evidence requiring longitudinal validation.

The argument that does not depend on resolving the causation contest is narrower and more durable. The cohort that would normally enter the workforce and accumulate the experience that produces the next generation of senior practitioners is being displaced before that accumulation begins. Senior software engineers, senior customer service managers, and senior operations specialists are not born senior. They are built from entry-level experience. When the entry-level positions disappear, the next generation of senior practitioners fails to form. The senior workforce of the decade beginning 2030 may not form at its projected scale, because the cohort that would build it is currently not building it.

PLATE II — THE LADDER PULLED UP.



SignalFire's 2025 State of Tech Talent Report⁹ documents a 50 percent drop in new graduate hiring at the 15 largest technology firms since 2019 and a 25 percent drop from 2023 to 2024 alone. The Stanford study measures relative employment decline in AI-exposed occupations across the full economy. The two data sources measure different phenomena at different levels of aggregation, and both point in the same direction. The training pipeline is not being interrupted at the surface; it is being interrupted at the foundation.


The terminus

The economic argument for replacement architecture proceeds from the firm's cost reduction and the investor's return on AI infrastructure investment. The argument stops at the quarterly or annual horizon. Extended to a multi-year population-scale horizon, the replacement architecture terminates in a social contract problem that some of its most prominent advocates have named without fully resolving a foundational objection to it.

Universal Basic Income has a serious policy literature. The Stockton Economic Empowerment Demonstration documented measurable improvements in employment, mental health, and financial stability in a randomized controlled design.¹⁰ The Finland Basic Income Experiment produced equivalent findings on wellbeing.¹⁰ Sam Altman's connection to the GiveDirectly East Africa basic income study¹⁰ is consistent with his repeated public positioning of UBI as the policy framework for an AI-displaced workforce. Andrew Yang's Freedom Dividend framing gave the policy a political vocabulary. None of this is the paper's critique.

The critique is structural and specific. UBI as a floor alongside meaningful work, where the income guarantee reduces the cost of transitions and supports capability development, is a policy instrument with established positive effects. UBI as a terminus, where the income guarantee is the dominant residual social contract after the workforce pipeline has collapsed and no augmentation path exists, is a different proposition. The terminus version is what the replacement architecture produces when operated at scale and sustained across decades. The paper distinguishes these two versions not to dismiss the floor but to name what the terminus means: an economy in which the route from capability development to productive contribution has been removed, and the remaining offer is income without participation. The scaffold is already built; the only variable left is whether the blade falls on a workforce that still knows how to climb or on a population trained only to receive.

PLATE III — THE SCAFFOLD ALREADY BUILT.



The foreclosing of the augmentation alternative is the architectural move that converts UBI from floor to terminus. Where augmentation deployment is operational, the income guarantee can complement work that builds capability. Where replacement deployment dominates, the income guarantee replaces work that would have built capability. The architecture, not the policy instrument, determines the terminus.

The evidence above documents Culture 1 as operational. It is not a projection or a warning about what may happen. Named executives at named organizations have named their architectural choice. The state apparatus of the most populous nation on earth has built the regulatory and infrastructural scaffolding to sustain it. The early-career cohort data shows the pipeline consequences at measurable scale. The lagging-indicator evidence from the most prominent replacement deployment shows the quality consequences at organizational scale. And the policy conversation about workforce income has already reached the terminus question.

The argument for a second architecture does not rest on the absence of Culture 1 consequences. It rests on the presence of an alternative that is also currently operational, also chosen by named executives at named institutions, and also producing documented outcomes. That architecture is examined next.

THE AUGMENTATION ARCHITECTURE

The live battleground

The United States is the contemporary site where both architectures are visible simultaneously and where named executives at named institutions are making opposite choices with the same frontier models in the same year. No architecture has been state-imposed; both coexist publicly. This is what makes the US the Culture 2 exemplar, not because the augmentation architecture is winning, but because the choice remains open and legible at scale.

The replacement architecture is currently outpacing the augmentation architecture in raw headcount-reduction count across the US technology sector. The cases documented in Arc 1 are US-headquartered or US-operating institutions. The augmentation architecture is operational at the largest US institutions measured by associate count, capital investment in workforce development, and named executive commitment. The choice is being made every quarter, in every deployment, by named people in named roles, and the contrast between those choices is currently visible in the quarterly filings, the analyst calls, and the statements executives make under conditions where reputational consequence follows.

Two US-headquartered institutions operating at maximum institutional scale have made the augmentation choice publicly, funded it materially, and produced measured outcomes against it.

JPMorgan Chase committed \$18 billion to technology investment in 2025, with AI as the central component. The bank expanded its AI-specialized workforce by 67 percent, from approximately 1,500 specialists in 2022 to a projected 2,500 in 2025, while simultaneously enrolling 300,000 workers in active training programs. The "AI Made Easy" internal training initiative reached tens of thousands of enrolled associates. Generative AI tools deployed across more than 200,000 employees across divisions, with the LLM Suite recognized as Innovation of the Year in the Americas category.

Documented analyst productivity outcomes included 2 to 4 hours of daily time savings on routine tasks, with that capacity redirected toward client-facing and judgment-intensive work rather than eliminated. Mary Erdoes, head of Asset and Wealth Management, made AI training mandatory for all new hires. Chief Analytics Officer Derek Waldron named the institutional mission directly: "Our mission is to make JPMorgan Chase a truly AI-connected enterprise." Jamie

Dimon's public framing of AI as something that will augment "almost every job" at the institution, with the added framing of "plenty of jobs" as the workforce outcome, represents the most prominent named-executive articulation of the augmentation architecture from the financial services sector.

Walmart's augmentation commitment operates at a different scale and in a workforce context that makes the architecture choice harder and the commitment more consequential. The \$1 billion commitment to skills training through 2026, combined with the OpenAI partnership launching free AI certification access across associates in 2026, extends augmentation infrastructure to 1.7 million US-and-Canada associates and 2.1 million globally. Walmart Academy has processed 3.5 million training participants, making it one of the largest private workforce development programs in operation globally. The proprietary Element machine learning platform deploys real-time translation, AI-assisted shift planning (reducing the planning time from 90 minutes to 30), and agentic workflow assistance across operations. Generative AI processing managed 850 million catalog data items, with the organization describing the scale as equivalent to 100 times the available human manpower for the task. CEO John Furner committed publicly in February 2026 to a position that two to five years out, Walmart expects "roughly the same number of people we have today." Chief People Officer Donna Morris offered the clearest statement of the augmentation choice as normative position, not just operational strategy: "We as big employers should be actively engaged in trying to equip our respective employees to be prepared for a world that is AI enabled and automated. It's unfortunate when companies use AI to replace workers instead of training them for what's ahead."

The Morris statement does not carry the weight of the Oracle/Ellison statement by default. What makes it load-bearing is the pairing: the two statements, Ellison's "We are choosing the chips" and Morris's "It's unfortunate when companies use AI to replace workers," are made by senior executives at two of the largest employers in the United States in the same year, about the same technology, and they represent opposite architectural commitments documented with capital allocation and operational deployment rather than aspiration.

One honest qualification is required and carried. Both JPMorgan and Walmart have conducted workforce reductions during the same period in which the augmentation commitments were made. The paper does not claim either institution as a pure Culture 2 deployment. The augmentation architecture holds at the level of stated policy, major program investment, quantified deployment outcomes, and named executive commitment. The multi-year longitudinal outcome data, which would measure whether the commitment produces the

compounding capability outcomes the architecture promises, does not yet exist. The programs are mid-deployment as of 2026. The paper carries the documented commitments and documented early outcomes as evidence of architectural choice in operation, not as proof of completed transformation.

The named architecture

Augmentation deployment treats AI as infrastructure for human capability rather than substitute for human labor. The deployment pattern produces a different sequence from the replacement architecture: structured cognitive demand on the human practitioner, AI-extended reach at the research and production layer, human judgment at every consequential decision point, measurable capability development in the practitioner over time, and workforce ladders that remain intact because the entry-level work, redefined as AI-governance work rather than AI-replaceable work, still exists and still builds toward senior competence.

The intellectual case against the augmentation architecture has a named formulation. Mo Gawdat, former Chief Business Officer at Google X, described across two interviews in March and April 2026 what he called the Fourth Inevitable, the fourth step in a sequence developed across his prior work: any organization or nation that develops superior AI capability will deploy it fully, and any organization or nation that does not match that deployment will become irrelevant to those that do. This framing treats the replacement architecture as determined by competitive pressure rather than chosen by named actors with available alternatives. The framing overreaches at the fourth step. It correctly describes what happens in the absence of governance infrastructure; it incorrectly treats that absence as proof that no governance infrastructure is possible. JPMorgan's 300,000 trained workers, Walmart's 1.7 million associates in scope, and the established productivity outcomes that accompany augmentation deployment are all produced under the same competitive pressure conditions that Gawdat describes as forcing replacement. The competitive pressure is real. The architectural choice within that pressure is not foreclosed.

The augmentation architecture is not a moral preference in competition with a strategic alternative. It is a principled position about where value accumulates across time. The replacement architecture captures short-horizon cost reduction. The augmentation architecture captures compounding capability, workforce ladders that remain climbable, and organizational knowledge that deepens rather than hollows. The economic case for the augmentation architecture is addressed in the next section.

The productivity and quality evidence

The Cynergy Bank case, introduced in the opening section, anchors the productivity-with-quality evidence at the individual-firm level. The 18 percent reduction in customer complaints and the 15 percent reduction in transactional call volume occurred because AI extended human relationship management capacity rather than substituting for it. The outcomes were quality outcomes and customer outcomes, not cost-reduction-only outcomes.

The research literature on augmentation-pattern deployment points in the same direction at scale. Noy and Zhang, publishing in *Science* in 2023,¹¹ documented a 40 percent reduction in task time and an 18 percent improvement in output quality in a randomized controlled design for knowledge workers using generative AI for writing tasks. Brynjolfsson, Li, and Raymond, publishing as an NBER working paper,¹² documented a 14 percent customer support productivity improvement in an enterprise deployment study. Both studies found something important beyond the headline productivity figure: the gains concentrated among workers who used AI to extend their own reasoning rather than to replace it. Workers who accepted AI outputs without verification produced lower-quality outcomes on tasks where the AI was wrong. The operative variable in both studies was the method of use, not the presence or absence of the tool.

The PwC 2025 Global AI Jobs Barometer¹³ documented a 56 percent wage premium for workers with AI skills compared to workers without them in the same occupational categories. Industries most exposed to AI produced four-fold productivity growth compared to industries least exposed over the period from 2022 through 2025. The wage premium is meaningful precisely because it documents human scarcity in augmented roles, not human elimination from them. Augmentation deployment creates a new form of human value: the capacity to govern AI effectively commands a wage premium because that capacity is currently scarce relative to demand for it.

The World Economic Forum Future of Jobs 2025 report¹⁴ projects 170 million roles created and 92 million roles displaced by 2030, a net gain of 78 million positions. Projection figures from aggregate forecasting institutions carry uncertainty that this paper does not minimize; the 2016 Oxford estimate of 47 percent of US occupations at high risk of automation has not materialized on that timeline, and aggregate projections have a consistent track record of underestimating augmentation adoption while overestimating displacement speed. The WEF figures are carried here not as settled forecasts but as documentation that the analytical institutions tracking this market project net positive role creation under augmentation scenarios, which is the direction the current operational evidence points.

The productivity evidence establishes that the augmentation architecture produces verifiable returns. The natural objection is whether those returns come at the cost of quality, and the Cynergy case, the Noy/Zhang study, and the Brynjolfsson enterprise data all show quality improving alongside or independently of productivity gains. The second natural objection is economic: does the augmentation architecture produce adequate returns to justify the investment? That question is addressed next.

The economic case

The augmentation architecture is not the morally correct alternative to the replacement architecture. It is the productively sound choice when measured across a business horizon long enough to capture compounding returns, workforce ladder sustainability, and quality outcomes.

McKinsey's generative AI economic analysis¹⁵ puts the annual value potential of generative AI at between \$2.6 trillion and \$4.4 trillion when the technology is deployed in augmentation rather than replacement patterns. The Penn Wharton Budget Model September 2025 analysis¹⁶ projected a 1.5 percent GDP increase by 2035 and \$400 billion in deficit reduction under AI-augmented workforce scenarios, with a permanent annual growth fade to below 0.04 percentage points acknowledged. These aggregate projections carry the same forward-facing uncertainty as the WEF jobs figures and are not presented as settled forecasts. They are carried as documentation that the economic analysis of augmentation deployment points toward material aggregate value.

The more immediate case is in the lagging-indicator evidence from replacement deployments that are now old enough to produce quality outcomes. The Klarna reversal is the most documented. Klarna announced in February 2024 that AI could handle the work of 700 customer service agents¹⁷ and implemented the corresponding operational model. By late 2025 and into 2026, CEO Sebastian Siemiatkowski described the outcome: "Cost, unfortunately, seems to have been a too predominant evaluation factor when organising this. What you end up having is lower quality." What was saved in dollars was lost in trust, and the company moved toward rehiring and hybrid human-AI models. The Gartner projection,¹⁸ published in the same period and consistent with early pattern evidence, holds that by 2027 half of companies that cut customer service staff because of AI will need to rehire due to quality deterioration. The prediction is not speculative at this point; it describes a pattern already visible in the leading cases.

The economic structure of the Klarna reversal is worth naming precisely. The cost reduction from replacement deployment is frontloaded and visible. The quality cost of replacement deployment is backloaded and invisible until the

quality deficit exceeds the tolerance threshold of customers, regulators, or both. Replacement deployment captures the short-horizon and forfeits the medium-horizon. Augmentation deployment requires investment in governance, measurement, and workforce development at the short-horizon and compounds returns at the medium-horizon. The accounting framework that makes replacement look profitable and augmentation look expensive is a short-horizon accounting framework applied to a medium-horizon outcome.

The financial stability argument for the augmentation architecture entered federal policy at the cabinet level in March 2026. US Treasury Secretary Scott Bessent, at the launch of the FSOC AI Innovation Series,¹⁹ stated that "failure to adopt productivity-enhancing technology is its own risk," treating augmentation delay as a stability threat rather than a safety preference. Deputy Assistant Secretary Christina Skinner reinforced the position: AI adoption is critical to America's financial stability and a precondition to economic growth. The FSOC's December 2025 Annual Report²⁰ had already established an AI working group charged with supporting AI adoption as a financial system resilience requirement. The Treasury framing converts the economic argument from corporate preference into a systemic risk frame: institutions that fail to develop governed AI capacity may become structurally less resilient than those that do, with consequences that extend beyond individual firm performance into the financial system as a whole.

The scalability limit and the honest concession

The paper has documented augmentation deployment at JPMorgan, at Walmart, and across the productivity research literature, and has made the economic case for the augmentation architecture's compounding returns. The scalability objection deserves a direct answer: is Culture 2 available only to organizations with patient capital, deep governance infrastructure, and the luxury of multi-year investment horizons?

This objection is not easily dismissed. The organizations producing the strongest augmentation evidence are among the largest and most capitalized institutions in their sectors. JPMorgan's \$18 billion technology investment is not available to a regional bank. Walmart's Walmart Academy and its 3.5 million training participants are not available to an independent retailer with 200 employees. The argument that Culture 2 is operationally available does not mean it is operationally equivalent for all organizations. It means the principles are operationally available, and the question of access to implementation is a separate question from the question of architectural availability.

The scalability proof the paper cannot yet make is the longitudinal multi-firm outcome study at the mid-tier enterprise level. The programs reviewed here

are mid-deployment as of 2026. Whether augmentation deployment at JPMorgan and Walmart scale produces the compounding capability outcomes the architecture promises across five to ten year horizons remains an open empirical question. That gap is carried explicitly in the limitations section.

What the paper can show is that the principles underlying the augmentation architecture are not capital-dependent in their own right. The core requirement for augmentation deployment is not a billion-dollar training budget. The requirement is governance: named human accountability at consequential decision points, measurement of whether practitioners are developing capability or offloading it, and incentive alignment between short-horizon cost metrics and medium-horizon capability metrics. Organizations of any size can choose to place a human with authority at every decision point where AI output becomes operational. Organizations of any size can measure whether their AI deployment is developing or eroding the cognitive capacity of the people using it. Organizations of any size can restructure accounting to capture medium-horizon value rather than optimizing exclusively for short-horizon cost reduction.

The access question and the principles question are separable, and the paper separates them deliberately. The principles are available. The access to large-scale implementation is not uniformly available. Both things can be true, and the paper carries both rather than resolving the tension prematurely.

The access question received a direct federal answer in March 2026. The US Department of Labor launched the Make America AI Ready initiative,²¹ delivering AI literacy through a seven-day SMS course accessible on any mobile phone, including basic devices without internet connection. Deputy Secretary of Labor Keith Sonderling described the goal as ensuring all Americans have the skills to share in AI-driven prosperity. The delivery choice signals where the federal government places the adoption gap: not in engineering depth but in foundational confidence and practical evaluation skill. The course maps directly to the Department of Labor's February 2026 AI Literacy Framework,²² which sequences learning through five content areas culminating in evaluation and responsible use, the same cognitive demand the augmentation architecture requires of practitioners at every deployment point. The capital requirement for that deployment is a text message.

The principles without the products

The governance architecture that distinguishes Culture 2 deployment from Culture 1 deployment is describable in principles without naming any specific implementation. This matters for a reason that is principled rather than rhetorical: the principles are testable by outcome, and any implementation of the

principles that produces the predicted outcomes satisfies the argument regardless of which organization built it or which framework it operates under.

Three principles characterize the deployment-side governance that separates augmentation from replacement at the operational layer.

The first principle is named human accountability at consequential decision points. A deployment in which AI produces output that flows directly into operational consequence without a named human exercising documented judgment at the boundary between AI output and operational effect is not augmentation deployment; it is replacement deployment with an optional human in the aesthetic position of reviewer. The named human must have the authority to stop the process, the responsibility to answer for the outcome, and the documentation that makes that authority and responsibility traceable. Where those three elements are present, governance is present. Where any of them is absent, governance is simulated.

The second principle is measurement of cognitive development rather than measurement of output volume. Augmentation deployment produces practitioners who are more capable over time than they were at the start of the deployment. Replacement deployment produces practitioners who are more dependent over time. The distinction is measurable, and the measurement gap, the absence of validated instruments for tracking human capability trajectory under AI collaboration, is one of the field's most significant current research needs. Instruments for this measurement are in development and available for validation, but the validation at population scale has not been completed. The paper carries the gap honestly and identifies the measurement problem as one requiring resolution before the augmentation architecture can be confirmed at the confidence level the replacement architecture's lagging-indicator evidence already supports.

The third principle is structural commitment survival under economic pressure. The Culture 2 architecture is not merely a program that exists when competitive pressure is low. The test of the architecture is whether the stated commitments survive the conditions, revenue shortfall, competitive threat, new market opportunity, that have historically caused voluntary commitments to fold. Culture 2 deployment that survives economic pressure at material commercial cost is different in kind from Culture 2 deployment that operates only when the cost of commitment is low. The empirical documentation of structural commitment survival is developed in the Mirror Diagnostic section that follows.

Any deployment architecture that satisfies all three principles, named accountability, cognitive development trajectory, and structural commitment survival under pressure, is what this paper means by Culture 2 deployment. The

body of the paper makes its case on these principles. One worked implementation, published by the author and disclosed in the appendix, is available for readers who want to study a specific operational form. The principles do not depend on the adoption of that implementation or any other specific framework. They describe what the architecture must do; the form in which any organization chooses to do it is a downstream decision.

With both architectures documented, the question becomes why a third architecture, one that forecloses replacement deployment through prohibition or sufficiently binding regulation, is not the superior alternative to the difficult internal choice between cultures. The Mirror Diagnostic begins the answer.

THE MIRROR DIAGNOSTIC

The argument

The distinction between Culture 1 and Culture 2 deployment is not a distinction between good and bad technology. It is a distinction between institutional choices made by named people under specific incentive structures. AI deployment patterns do not arrive from outside an institution's existing commitments, competitive pressures, and accounting frameworks. They inherit them. The technology is a mirror. What it reflects is the deployment decision made by the institution that purchased it, and that decision is made at the layer where institutional incentives operate, which is the deployment layer, not the capability layer.

This has a direct implication for where governance must work. Capability-side governance, whether regulatory, self-imposed by frontier labs, or advocated by safety researchers, addresses the development of AI systems. Development governance can slow, redirect, or shape what gets built. It cannot determine how what gets built gets deployed, because the deployment decision is made downstream by institutions with different incentive structures, different accountability mechanisms, and different time horizons from the organizations that built the model. A frontier lab that commits to safety at the development layer still sells its products to enterprises optimizing for quarterly returns. The commitments made during development do not travel into the deployment context unless deployment-side governance carries them there.

This does not make capability-side governance worthless. It makes capability-side governance incomplete as a solution to the Culture 1 problem. The paper's argument is not that development-layer commitments are irrelevant; it is that they address the wrong layer when the architectural choice between cultures is being made. Two institutions deploying the same frontier model can produce Culture 1 and Culture 2 outcomes simultaneously, as the Oracle and Cynergy cases document. The model is the same. The architecture is different. The architecture is determined at deployment.

Hinton: the diagnostician

Geoffrey Hinton occupies a specific and legitimate position in this argument, as the architect who helped birth the revolution and now warns it may consume its creators. His CNN interview of December 2025 stated that 2026 would see AI "replace many, many jobs."²³ His estimate of a 10 to 20 percent probability of

extinction-level risk from artificial general intelligence, published in the Guardian and cited in CGTN in December 2024,²⁴ positioned him as the most prominent credentialed voice on the risk side of the capability debate. His April 2026 address in Geneva described AI development as a "fast car with no steering wheel,"²⁵ one of the most compressed formulations of the current governance deficit yet offered by a figure of his institutional standing. His exchange with Senator Bernie Sanders, in which he noted that tech billionaires had not thought through the consequence that "if workers don't get paid, there's no one to buy their products," added economic critique to what had previously been framed primarily as a safety argument.

Hinton's diagnosis identifies a real risk. The prescriptions associated with his position aren't complete, and the paper names that incompleteness without dismissing the diagnosis. His primary prescription has been some form of pause or international agreement to slow capability development. The foundational problem with the pause prescription is documented in Arc 3c. The narrower problem named here is that even if the pause prescription succeeded, it would address the development layer while leaving the deployment layer ungoverned. Organizations operating under paused-development conditions could still make Culture 1 deployment choices with systems already built. The culture produced by the deployment is not determined by the development timeline; it is determined by the deployment decision.

Yann LeCun has contested Hinton's extinction probability estimate, describing the scenario as technically implausible under current architectural assumptions. That dissent is preserved here as required. The paper doesn't resolve the Hinton/LeCun disagreement on extinction risk because the paper's argument does not depend on resolving it. Whether or not extinction risk is in the plausible range, the Culture 1 deployment pattern documented in Arc 1 is operational now, producing measurable consequences now, and addressable through deployment-side governance now. The magnitude of the long-run risk is a separate question from the availability of the near-term architectural solution.

The Anthropic sequence

The most precisely recorded case of deployment-incentive inheritance in the current record is the sequence of decisions made by Anthropic between February 9 and April 8, 2026. The sequence requires complete narration here because a simplified version of it, which several analytical platforms produced when given only partial prompts, loses the argument's full pattern. The argument depends on the full pattern: a voluntary commitment that folded, a public redline commitment that survived, the same firm, the same pressure cycle, and opposite

outcomes determined by the architecture of the commitment rather than the goodwill of the organization.

On February 9, 2026, Mrinank Sharma, who led Anthropic's Safeguards Research team and was responsible for the alignment work undergirding the company's safety commitments, submitted his resignation. His letter, reported by Forbes, Yahoo Finance, and CNN,²⁶ stated: "The world is in peril." Its significance is institutional: the person whose professional function was to maintain Anthropic's internal alignment between stated safety commitments and operational decisions concluded that the commitment architecture was not holding.

PLATE IV — THE REFUSAL AT THE GATE.



On February 24, 2026, Defense Secretary Pete Hegseth met with Dario Amodei and presented a demand: grant the Pentagon unrestricted "all lawful use" access to Claude by 5:01pm Eastern on February 27, or face Defense Production Act invocation and supply chain risk designation. The \$200 million Pentagon contract signed in July 2025 was the explicit commercial stake. The demand was not a negotiating position; it was a deadline with named consequences.

On the same day, February 24, Anthropic published Responsible Scaling Policy version 3.0. The new policy removed the binding pause commitment that had been a centerpiece of the prior version. Internal discussions on the change had reportedly run for nearly a year. The board approved the removal unanimously. The Responsible Scaling Policy had been the primary public evidence that Anthropic's safety commitment was not merely aspirational but operational and enforceable. Its pause provision was what distinguished Anthropic's safety positioning from competitors who expressed safety concern without committing to specific operational constraints. The pause provision was also, by the characterization of GovAI analysts, voluntary and self-assessed: Anthropic would determine when its own systems met the criteria that would trigger a pause, with no external verification mechanism.

On February 26, Anthropic refused the Pentagon's demands regarding autonomous weapons development and mass domestic surveillance. Dario Amodei's statement was specific:²⁸ the company "cannot in good conscience accede" to those requirements. The refusal was public, on record, and consequential. The autonomous weapons and surveillance commitments were structural in the relevant sense: they were public redlines with enough legal specificity to be legally defensible, recorded in a form that external parties could reference, and attached to named accountability at the CEO level.

On February 27, the Trump administration directed all federal agencies to cease using Anthropic products. Pete Hegseth designated Anthropic a supply chain risk.²⁹ The \$200 million contract and the broader federal contracting relationship were effectively terminated. Anthropic's public redline commitment had survived the test at material commercial cost.

On March 9, Anthropic filed suit against the federal government in the Northern District of California (Case No. 3:26-cv-01996).³⁰ On March 26, Judge Rita Lin issued a preliminary injunction blocking the agency ban, with a 43-page ruling finding that the retaliatory actions were likely illegal (Document 134, N.D. Cal.).³¹ On April 8, the DC Circuit denied Anthropic's request for a stay pending review (Case No. 26-1049).³²

The full sequence is the evidence. Jared Kaplan, Anthropic's chief science officer, stated in a TIME interview published February 24, 2026:³³ "I don't think we're making any kind of U-turn." The record complicates that characterization: the pause commitment dropped, the weapons and surveillance commitments held, and the distinction between them is the architecture of the commitment, not the sincerity of the commitment-makers.

The pattern named

Anthropic's own policy text states the logic that produced the RSP pause-commitment removal. The Responsible Scaling Policy v3.0²⁷ explains the decision with the following reasoning: "If one AI developer paused development to implement safety measures while others moved forward training and deploying AI systems without strong mitigations, that could result in a world that is less safe — the developers with the weakest protections would set the pace, and responsible developers would lose their ability to do safety research and advance the public benefit." This is the precise dynamic at work: the tendency for deployment incentives to override governance regardless of institutional intent, documented here in Anthropic's own policy reasoning. Anthropic is not being criticized for naming this dynamic accurately. Anthropic is being cited as primary-source evidence that the organizations most committed to safety have themselves identified the systemic reason why voluntary capability-side commitments fold: the competitive structure penalizes unilateral restraint while the deployment consequences of that restraint accrue to the broader public rather than to the organization bearing the cost.

The distinction between the RSP pause commitment and the weapons and surveillance commitments is the operational lesson. The pause commitment was self-assessed, self-enforced, and had no external verification mechanism. Its removal cost Anthropic nothing observable in the short term. The weapons and surveillance commitments were public, specific, legally defensible redlines attached to named accountability at the CEO level. Their maintenance cost Anthropic \$200 million in contract revenue and produced a federal supply-chain-risk designation. The public redline commitments survived the test precisely because their architecture made the cost of abandonment public, named, and legally consequential. The voluntary commitment folded precisely because its architecture made the cost of abandonment invisible and non-consequential.

This is the deployment-side governance principle in empirical form. It is not a theoretical claim about what governance should look like. It is a documented case in which one commitment architecture survived and another did not, at the same institution, under the same pressure, in the same week.

THE GOVERNANCE ARCHITECTURE

What structural commitment requires

The Anthropic case documents the principle in the negative: a commitment without external verification, named accountability, and legal defensibility folds under competitive pressure. The principle in the positive is that deployment-side governance with those three properties produces commitments that survive.

The capability-side governance debate, which occupies most of the current AI safety field's intellectual resources, concentrates on what AI systems should be allowed to do. The deployment-side governance question is different: who is accountable when a named AI system produces a named consequence for a named person, and what accountability channels exist for that person to exercise? The capability-side question matters, but the deployment-side question is the one that determines the cultural outcome of the choice being made in every enterprise deployment today.

Three tiers of governance exist in current practice, and only one of them produces the structural commitment architecture that survived the Anthropic February pressure test.

The first tier encompasses principled statements without enforcement mechanisms. Organizational AI ethics documents, responsible use policies, and voluntary commitments expressed in public communication without operational binding are at this tier. The statements are real. The accountability mechanism is absent. When competitive pressure arrives, the cost of abandoning a first-tier commitment is the reputational consequence of the abandonment, and only if the abandonment is visible. The RSP pause commitment operated near this tier: public enough to be reputationally consequential, but voluntary and self-assessed enough that the cost of removal could be managed through messaging rather than institutional consequence.

The second tier encompasses machine checks machine: automated compliance monitoring, algorithmic output flagging, and AI safety systems that evaluate AI output against specified criteria. Second-tier governance is a genuine advance over first-tier governance; it catches classes of failure that first-tier governance misses entirely. Its ceiling is the absence of a named human who bears personal accountability for the outcome. When an automated safety system flags an output and no named person is required to document a judgment call about what to do with the flag, the flag becomes data rather than governance.

Second-tier governance can generate audit trails without generating accountability. No named human answers personally when the second-tier system fails.

The third tier requires a named human being with binding checkpoint authority at every consequential AI decision point, connected to at least one of four accountability channels: moral accountability (the person is on record with a named judgment call), employment accountability (the person's role depends on the quality of that judgment over time), civil accountability (the person is legally exposed if the judgment produces identifiable harm to identifiable parties), and criminal accountability (the person faces legal jeopardy if the judgment produces harm meeting criminal thresholds). The third tier is not designed for every AI output; it is designed for every AI output that becomes operational consequence. The distinction matters: AI output that reaches a customer, affects an employment decision, influences a clinical outcome, or shapes a credit determination has crossed the boundary from output to consequence. At that boundary, the third tier is what the governance architecture requires. This is the checkpoint that turned Anthropic's public redline into survival while the pause commitment folded.

PLATE V — THE CHECKPOINT.



The checkpoint principle

The checkpoint principle, as an operational concept, holds that named human accountability at the boundary between AI output and operational consequence is the structural requirement that distinguishes Culture 2 governance from Culture 1 or Culture 1.5 governance. Culture 1.5 is the common failure mode: organizations that believe they are operating Culture 2 architecture because they have humans in the aesthetic position of reviewer, present in the process but not exercising traceable judgment with accountable recordkeeping. The Anthropic February sequence provides the test that distinguishes Culture 2 from Culture 1.5: does the governance commitment survive competitive pressure at material commercial cost? If the human review process is the first casualty of a cost-reduction initiative, the review process was not a checkpoint; it was an interface element.

The checkpoint principle is portable across organizational size, sector, and technology vendor. It does not require a specific software implementation, a specific framework, or a specific vendor relationship. It requires answering three questions before any AI output reaches operational consequence: who is the named human making the judgment call about this output; what record captures that judgment call in a form that is retrievable and auditable; and what accountability channels connect the named human to the consequence of the call. Where all three are present, the checkpoint is structural. Where any one is absent, the checkpoint is cosmetic.

Deployment-side governance in operational form

The principles above are describable in terms of what any deployment-side governance architecture must produce, without naming any specific implementation. Five operational requirements characterize deployment-side governance that satisfies the Culture 2 architecture.

The first requirement is structured cognitive demand. The deployment must require the human practitioner to perform cognitive work that builds expertise rather than outsourcing cognitive work that would have built expertise to the AI system. This does not mean humans should perform tasks that AI does more efficiently. It means the deployment must be designed so that the AI extends the human's reasoning capacity rather than replacing it. The test is measurable: does the practitioner who has worked within this deployment architecture for twelve months show more or less capacity for independent judgment in AI-adjacent tasks than the practitioner who began the same work twelve months prior?

The second requirement is dissent documentation. In any deployment using multiple AI sources, parallel AI outputs, or AI-mediated research, the governance architecture must record cases where sources or outputs disagree and preserve

that disagreement for human review rather than resolving it algorithmically. Dissent resolution is a judgment call. Judgment calls are the sites of accountability. Algorithmic dissent resolution produces outputs without producing accountability.

The third requirement is calibrated confidence with conditional knowledge. Outputs that reach operational consequence must carry a documented confidence level and an explicit statement of the conditions under which the confidence level would change. An output that is correct under current conditions but incorrect if a specified assumption changes is not an unconditional output. Treating it as unconditional is a governance failure at the output layer before the deployment layer is even reached.

The fourth requirement is organizational architecture for capability formation. The workforce must be structured so that AI-adjacent roles remain developmental. Entry-level positions must exist in forms that build toward senior competence rather than forms that the AI system has absorbed. The training ladder must remain climbable. This requirement is organizational design, not technology design, and it is exactly the requirement that the replacement architecture violates at scale.

The fifth requirement is cognitive amplification measurement over time. The organization must measure whether the deployment architecture is producing practitioners who are more capable, less capable, or unchanged in their capacity for AI-independent judgment over time. Without this measurement, the organization cannot distinguish Culture 2 deployment from Culture 1 deployment that has been successfully branded as Culture 2. The measurement instrument for this distinction is an active research priority, not a solved problem, and the paper carries that gap honestly.

One worked implementation of these principles exists in open-source published form, developed and maintained by this paper's author, available at the repository referenced in the appendix. That implementation is referenced in the appendix as exhibit for readers who want to study a specific operational form. The principles do not depend on that implementation and do not require its adoption. Any implementation producing measurable cognitive development, documented dissent handling, calibrated outputs, intact career ladders, and longitudinal capability measurement satisfies the Culture 2 architecture as defined by this paper.

THE ABSENT CULTURE

The institutionally serious attempt

The argument that a third architecture, one that forecloses the Culture 1 choice through sufficiently binding prohibition or regulation, is not available as a durable option does not rest on the claim that no serious institution has tried to build it. The European Union has tried. The EU AI Act is the contemporary world's most institutionally serious attempt to construct Culture 3, and what the EU AI Act documents, through its implementation challenges rather than through its political ambitions, is exactly what the paper argues: even maximum institutional and political will cannot make the prohibition architecture coalesce against globally distributed capability operating on commercial timelines.

The EU AI Act (Regulation 2024/1689),³⁴ adopted in 2024, established the world's first comprehensive AI risk-tier framework across four categories: unacceptable risk, high risk, limited risk, and minimal risk. Penalties reach up to €35 million or 7 percent of global annual turnover, a regime exceeding GDPR in financial severity. Extraterritorial scope extends the act to any AI system whose outputs reach persons located in the EU, regardless of where the provider or developer is headquartered. The regulatory architecture is serious, penalty-backed, and globally reaching in its jurisdictional claim.

The phased implementation timeline set the Act's prohibited practices and AI literacy obligations to apply from February 2, 2025. General purpose AI model obligations followed from August 2, 2025, with the GPAI Code of Practice published July 10, 2025, operating as a voluntary safe harbor carrying presumption of conformity. High-risk AI obligations were set to apply from August 2, 2026. The structure was ambitious and the timeline was statutory.

The friction materialized before the statutory deadlines arrived. The European Commission missed its own internal deadline for issuing Article 6 classification guidance, due February 2026 and not delivered on schedule. Harmonised technical standards, which organizations need to meet conformity with high-risk obligations, accumulated delays across standards bodies. Notified body capacity, the assessment organizations authorized to certify AI system conformity, remained insufficient relative to projected demand across member states. Organizations building compliance programs against the August 2026 high-risk deadline were doing so against standards that had not fully

materialized. The conformity architecture was being demanded before the conformity architecture was available to be built.

Industry pressure followed. The EU AI Champions Initiative, in an open letter signed by executives at Airbus, ASML, Lufthansa, Mistral, and TotalEnergies, requested a two-year clock-stop. A separate letter from ASML, Airbus, Ericsson, Nokia, SAP, Siemens, and Mistral AI named the risk directly: European organizations could face a competitive disadvantage relative to US and Chinese counterparts not subject to the same conformity burden, in the same window that AI capability was most consequential to competitive position. The Commission's initial public position, stated July 4, 2025, held firm to the statutory deadlines.

By May 7, 2026, EU co-legislators reached political agreement on the Digital Omnibus package,³⁵ extending the high-risk AI obligation deadline from August 2, 2026 to December 2, 2027 for Annex III standalone systems and to August 2, 2028 for AI embedded in regulated products. A sixteen-month delay on the Act's most substantive provisions, achieved not by a change in political will but by the inability to assemble the conformity infrastructure against globally operating commercial capability on statutory timelines.

The EU is not failing morally. The regulatory diagnosis motivating the Act is accurate: AI systems affect employment decisions, creditworthiness assessments, law enforcement interactions, healthcare triage, and educational admissions. These are consequential applications requiring accountability, transparency, and redress. The EU AI Act's risk-tier framework identifies that requirement correctly. The prescription, regulation-as-prohibition through binding extraterritorial conformity assessment, does not stabilize at deployment scale against capability that is globally distributed, commercially motivated, and operating on timelines faster than the standards bodies can produce the conformity architecture. The EU shows the implementation burden of prohibition architecture under maximum institutional effort, and that burden is the strongest evidence the third option does not coalesce at deployment scale.

One additional distinction is required for the argument to be precise. The EU AI Act, at its deployment-facing provisions, is not exclusively prohibition-regulation. Its transparency and human oversight requirements for high-risk applications are deployment-side governance in the paper's terms: they require named human accountability at specific decision points. Those provisions are consistent with the Culture 2 architecture. The EU AI Act's difficulty is that it pursues both prohibition-regulation and deployment-governance simultaneously, and the prohibition-regulation component is what the conformity architecture cannot assemble. The deployment-governance component is structurally

coherent and represents exactly the kind of accountability-regulation the paper advocates. The implementation delay damages both components, but the structural argument applies to the prohibition component specifically.

The pattern

The EU AI Act is the contemporary instance of a recurring pattern that has repeated across three prior capability domains.

Alcohol prohibition in the United States (1920 to 1933) identified a real harm, documented the harm credibly, and produced binding national legislation with criminal penalties. At peak enforcement in 1925, an estimated 30,000 speakeasies operated in New York City alone. By Prohibition's end in 1933, adulterated spirits produced by unregulated manufacturers had caused an estimated 10,000 deaths, with the mortality accelerating after 1927 when the federal government doubled the poison content of industrial alcohol to discourage bootleggers. The harm the prohibition movement identified did not diminish under prohibition; it transformed into a different and less governable form. The repeal in 1933 came not from the failure of the original diagnosis but from the documented failure of the prescription: the capability existed, the economic incentive existed, the demand existed, and the enforcement architecture could not match the combined force of all three.

The international drug control framework, formalized through the Single Convention on Narcotic Drugs in 1961 and expanded through subsequent decades, has operated for over sixty years against drug markets that have grown rather than contracted over that period. Survey and public health evidence across drug control regimes shows that demand consistently exceeds legal supply and generates illicit supply chains that adapt faster than the enforcement architecture can respond. The capability to produce and distribute controlled substances is globally distributed, economically motivated, and cannot be recalled by enforcement. The diagnosis motivating drug prohibition is accurate. The prescription does not stabilize against these conditions.

The Crypto Wars of the 1990s documented the same pattern at a shorter timescale. The US government's Clipper Chip proposal would have required a government backdoor in all encryption implementations, enabling lawful intercept. The proposal was technically defeated when Matt Blaze, Whitfield Diffie, and other researchers showed that the architecture weakened security against all adversaries, not only law enforcement. But the argument that ultimately closed the Crypto Wars debate was capability distribution: by the mid-1990s, strong encryption implementations were available globally, open-source, and beyond any government's capacity to recall. Phil Zimmermann had already

published PGP. The capability was distributed. Prohibition was no longer a viable option regardless of legislative will.

Three failure modes recur across all four cases: global coordination breakdown under economic pressure, the permanent enforcement requirement that capability-based prohibition imposes once the capability exists and cannot be recalled, and the decentralization of capability that makes enforcement increasingly expensive relative to circumvention as distribution expands. AI training capability is currently available on consumer GPUs. Open-weights model releases, multiple of them from frontier labs, have distributed foundational AI capability to any operator with hardware access and the motivation to use it. The international coordination mechanism that would be required to enforce global capability prohibition does not exist and has not been successfully assembled in analogous domains in the prior century.

STOPAI and PauseAI

STOPAI and PauseAI occupy the civil society expression of the Culture 3 prescription. Both identify real concerns, and both have publicly tried to convert those concerns into political pressure. The difficulty is not sincerity. The difficulty is institutional coherence. Prohibition movements require broad legitimacy, stable organizational discipline, and enforceable policy mechanisms. The current anti-AI movement has not yet assembled those conditions at the scale required to become a durable governance alternative. In 2026, public conflicts and reported incidents around the anti-AI movement showed how difficult it is for prohibition-oriented activism to maintain broad legitimacy, message discipline, and policy coherence under public pressure. The movement-level evidence points in the same direction as the institutional evidence: the prohibition prescription does not coalesce.

The choice

Two architectures are operationally available. One is documented in layoff data, cognitive offloading evidence, training-pipeline collapse, and UBI as the only remaining social contract at its terminus. The other is documented in productivity research, wage premium data, projected net job creation, named executive commitments at the two largest US employers by associate count, and published infrastructure for the governance architecture that distinguishes augmentation from replacement. The third option is not a structural option. It has been attempted at maximum institutional scale, it has failed for structural reasons that recur historically, and the capability distribution that closes it has already occurred.

The future of work under AI is not a forecast. It is a choice. The choice is being made every quarter, in every deployment decision, by named executives at named institutions. The paper has documented both operational architectures, established that both are chosen today by real organizations with real capital and real named accountability, documented why the prohibition architecture does not coalesce against globally distributed commercially motivated capability, and identified the governance principles that determine which architecture an institution actually operates regardless of which architecture it claims to operate.

The choice is not this paper's to make. It is yours.

LIMITATIONS

Six limitations constrain the argument and are disclosed here rather than carried silently.

Culture 2 scalability at enterprise scale is not yet proven. The paper's primary Culture 2 deployment evidence comes from JPMorgan Chase and Walmart, two of the largest and most capitalized institutions in their respective sectors, and from Cynergy Bank, a single mid-size financial services firm. PwC and WEF figures are aggregate forecasts with the forward-looking uncertainty those methods carry. The paper claims Culture 2 is operationally available; it does not claim that Culture 2 dominates aggregate deployment outcomes or that the organizational conditions for augmentation deployment are uniformly accessible. The scalability proof requires longitudinal multi-firm outcome data at the mid-tier enterprise level that does not yet exist.

The primary evidence gap is longitudinal. Validation of structured human-AI collaboration as a sustained enterprise deployment model, measured across multiple organizations over multi-year horizons under competitive pressure, is the research need this paper cannot satisfy. A cognitive amplification measurement instrument designed to close this gap exists in the research literature and is documented in the appendix. Its longitudinal validation at population scale is a current research priority, not a completed result.

The AI-washing critique is acknowledged and applied. Not all AI-attributed layoffs are AI-caused. The paper applies strict three-tier attribution discipline: Tier A requires named-executive primary attribution; Tier B covers AI-adjacent restructuring with preserved dissent from leadership denials; Tier C excludes independently denied cases from the primary evidence base. The aggregate strictly-AI-caused layoff count is smaller than the aggregate AI-attributed count. The Tier A evidence base remains substantial after the attribution discipline is applied.

The cognitive evidence is mechanistically suggestive, not longitudinally tested. The MIT Media Lab EEG study, the Gerlich correlational survey, and the Macnamara theoretical analysis all point in the same direction. None of them establishes the causal claim, that unstructured AI-dependent workflows produce persistent measurable reductions in cognitive capability, at the standards adjacent research fields require for population-level intervention claims. The paper applies the epistemic discipline established in the author's prior published critique of the cognitive decline literature, does not

overclaim in the direction of the finding, and preserves the opposite claim's absence of support as an equally relevant statement about the current evidence base.

The author corpus dependency is acknowledged. The Culture 2 prescription in this paper draws on research, frameworks, and case studies published by the author. The appendix carries the full disclosure. The underlying argument that Culture 2 deployment is available and distinguishable from Culture 1 deployment by specific operational characteristics does not depend on adoption of the author's implementations; those implementations are one available expression of principles that any organization can satisfy through its own architecture. The paper does not claim the author's implementations are superior to alternatives that satisfy the same principles. Independent implementations that produce measurable cognitive development, documented dissent handling, and governance commitment survival are equally valid expressions of the Culture 2 architecture.

Several current-event claims require source lock before publication. The paper relies on rapidly developing 2025 and 2026 events involving corporate restructurings, AI policy disputes, litigation, and regulatory implementation. Claims involving direct executive quotations, court proceedings, and exact workforce figures should receive a high volatility source refresh immediately before final publication, as active litigation and corporate reporting may update. The Oracle restructuring announcement, the Anthropic litigation sequence, and the EU Digital Omnibus implementation dates represent the highest verification priority.

The Hinton diagnosis is accepted; the prescription is not endorsed as complete. The paper does not contest Geoffrey Hinton's diagnosis of capability-side risk or workforce displacement risk. It locates the gap between an accurate diagnosis and a deployable prescription, and it does not resolve the Hinton/LeCun disagreement on extinction probability because the paper's argument does not depend on resolving it.

CLOSING

Larry Ellison chose the chips. Cynergy Bank chose the people who use them better. Both decisions were made in 2026. Both used frontier model technology. One funded its infrastructure commitment by removing thousands of people from its payroll. The other funded its quality commitment by ensuring the people remained.



The paper has documented what each architecture looks like in operation, what it produces across the evidence base available now, and where that evidence base is strong and where it is thin. It has shown why the third option, the option that forecloses the replacement choice through prohibition or sufficiently binding regulation, does not coalesce for foundational reasons that repeat across a century of analogous capability domains. And it has described the governance principles that determine which architecture an institution actually operates, regardless of which architecture it claims to operate.

Two architectures are available. The choice between them is consequential. It is being made every quarter, in every deployment, by named people in named roles at named institutions. It will continue to be made every quarter after this paper is published, in organizations this paper will never reach, by people who have never encountered its argument.

The architecture that compounds human capability requires governance. The governance requires named accountability. The accountability requires an organizational decision to build it. That decision belongs to you.

This paper borrows its organizing structure from Charles Dickens's *A Tale of Two Cities*, published in weekly installments in *All the Year Round* beginning April 30, 1859. The opening anaphora, the binary framing of two cultures operating simultaneously on the same historical stage, the insistence that the choice between them belongs to living people rather than to historical forces, and the use of named institutions as the narrative anchors rather than invented characters — all of these owe a direct debt to Dickens's method. The debt is acknowledged with gratitude. The argument is the author's own.

The principles described in this paper have been articulated by multiple researchers and practitioners across multiple implementations. The author of this paper has published one such implementation, documented in the appendix that follows. The appendix is an exhibit, not a prescription. The structural argument stands on its own terms.

APPENDIX: ONE IMPLEMENTATION OF THE PRINCIPLES

The body of this paper makes its case on principles. Many implementations of these principles exist or are possible. The author has published one, documented below for readers interested in studying a worked example. The appendix is an exhibit, not a prescription. Readers who find the principles compelling are not bound to this implementation; readers who find this implementation useful are not bound by it either. The structural argument carries the paper. The appendix carries the disclosure.

Principle from the body	One author's implementation (one of multiple available)
Deployment-side governance with named human accountability at consequential decision points; binding checkpoints with four accountability channels: moral, employment, civil, criminal	Checkpoint-Based Governance (CBG v5.0), open-source publication, basilpuglisi.com
Structured cognitive demand: role assignment, evidence evaluation, conflict documentation, calibrated confidence, explicit human judgment at every consequential output boundary	HAIA-RECCLIN Third Edition, open-source publication, basilpuglisi.com
Multi-AI orchestration with mandatory dissent preservation across parallel platforms; source-authority discrimination; synthesizer audit with overlap mapping	HAIA-CAIPR Specification v1.2.4, open-source publication, basilpuglisi.com
Governance orchestration policies with deterministic audit trail, hash-chained logging, and enforcement at the deployment layer	GOPEL Canonical Public v1.5, open-source publication, basilpuglisi.com
Organizational architecture for capability formation: trust and transparency, rhythm and culture, outcome anchoring	Growth OS Parts 1 and 2, basilpuglisi.com
Cognitive amplification measurement; discriminating offloaded thinking from persistent cognitive development; Human Enhancement Quotient / Augmented Intelligence Score	HEQ/AIS Scoring Rubric Specification v1, open-source publication, basilpuglisi.com and SSRN Abstract ID 6583419
Three-tier governance distinction: principles without enforcement (Ethical AI) / machine checks machine (Responsible AI) / named human accountability with binding checkpoints (AI Governance)	Capstone Reference v3.7, basilpuglisi.com; also Checkpoint-Based Governance v5.0

The implementation listed above is one expression of the principles argued in this paper. It is not the test of the principles. The principles are tested by whether they describe a deployment architecture that produces measurable cognitive development, governance commitment survival, and named human accountability. Any implementation that produces those outcomes satisfies the principles. The author's contribution is one worked example. Other contributions are welcome.

REFERENCES

- [1] Oracle restructuring and executive statements: Oracle Corporation. (2026, March 10). *Form 10-Q for the quarterly period ended February 28, 2026* (Fiscal 2026 Restructuring Plan). U.S. Securities and Exchange Commission. Reuters. (2026, March 31). Oracle begins layoffs affecting thousands, CNBC reports. <https://www.reuters.com/business/world-at-work/oracle-begins-cutting-thousands-jobs-cnbc-reports-2026-03-31/> | The Ellison and Catz formulations appear in secondary financial press (CX Today, Tech Insider, HR Reporter) attributed to a Stargate investor briefing on April 2, 2026. No primary Oracle transcript was available at review.
- [2] Google Cloud. (n.d.). *Cynergy Bank: Transforming customer service with AI-powered digital banking*. Google Cloud Customer Stories. <https://cloud.google.com/customers/cynergy-bank>
- [3] Carnegie Endowment for International Peace. (2025, July). *Crossroads era: China's AI governance and the post-DeepSeek landscape*. Carnegie Endowment for International Peace.
- [4] Challenger, Gray & Christmas. (2025, December). *Annual job cuts report 2025*. <https://www.challengergray.com> | Challenger, Gray & Christmas. (2026, April). *April 2026 job cuts report*. <https://www.challengergray.com>
- [5] Skillsyncer. (2026). *2026 tech layoffs tracker* [Live data dashboard]. <https://skillsyncer.com/layoffs-tracker>
- [6] Kosmyrna, N., Hauptmann, E., Tuan, Y. T., Olwal, A., & Maes, P. (2025). *Your brain on ChatGPT: Accumulation of cognitive debt when using an AI assistant for essay writing task*. arXiv preprint arXiv:2506.08872. <https://arxiv.org/abs/2506.08872>
- [7] Gerlich, M. (2025). AI tools in society: Impacts on cognitive offloading and the future of critical thinking. *Societies*, 15(1), 6. <https://doi.org/10.3390/soc15010006>
- [8] Brynjolfsson, E., Chandar, B., & Chen, R. (2025, August). *Canaries in the coal mine? Six facts about the recent employment effects of artificial intelligence* (Working Paper). Stanford Digital Economy Lab. Revised note published March 2026.
- [9] SignalFire. (2025). *State of tech talent report 2025*. <https://signalfire.com/state-of-tech-talent/>
- [10] Stockton Economic Empowerment Demonstration (SEED). (2021). *SEED findings: A guaranteed income pilot*. <https://www.stocktondemonstration.org> | Kangas, O., Jauhiainen, S., Simanainen, M., & Ylikännö, M. (2021). *Experimenting with unconditional basic income: Lessons from the Finnish basic income experiment 2017–2018*. Edward Elgar Publishing. | GiveDirectly. (2023). *GIVEDirectly's basic income program: Research and outcomes*. <https://www.givedirectly.org/research-at-givedirectly/>
- [11] Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187–192. <https://doi.org/10.1126/science.adh2586>
- [12] Brynjolfsson, E., Li, D., & Raymond, L. R. (2023). *Generative AI at work* (NBER Working Paper No. 31161). National Bureau of Economic Research. <https://doi.org/10.3386/w31161>
- [13] PwC. (2025). *2025 Global AI jobs barometer: AI skills now command 56% wage premium*. PwC Global. <https://www.pwc.com/gx/en/services/ai/ai-jobs-barometer.html>
- [14] World Economic Forum. (2025, January). *Future of jobs report 2025*. World Economic Forum. <https://www.weforum.org/reports/future-of-jobs-report-2025>
- [15] McKinsey & Company. (2023, June). *The economic potential of generative AI: The next productivity frontier*. McKinsey Global Institute.

- <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-ai>
- [16] Penn Wharton Budget Model. (2025, September 8). *The projected impact of generative AI on future productivity growth*. University of Pennsylvania. <https://budgetmodel.wharton.upenn.edu/>
- [17] Siemiatkowski, S. (2024, February). Klarna AI deployment announcement [Press release]. Klarna. Siemiatkowski, S. (2026). Quality assessment statement. Reported in Forbes and Fast Company.
- [18] Gartner. (2026, February 3). *Gartner predicts half of companies that cut customer service staff due to AI will rehire by 2027* [Press release]. <https://www.gartner.com/en/newsroom/press-releases/2026-02-03-gartner-predicts-half-of-companies-that-cut-customer-service-staff-due-to-ai-will-rehire-by-2027>
- [19] U.S. Department of the Treasury. (2026, March 20). *Treasury launches the Artificial Intelligence (AI) Innovation Series* [Press release sb0421]. <https://home.treasury.gov/news/press-releases/sb0421>
- [20] Financial Stability Oversight Council. (2025, December 11). *FSOC 2025 annual report*. U.S. Department of the Treasury. <https://home.treasury.gov/system/files/261/FSOC2025AnnualReport.pdf>
- [21] U.S. Department of Labor. (2026, March 24). *U.S. Department of Labor launches "Make America AI Ready" initiative* [Press release]. <https://www.dol.gov/newsroom/releases/osec/osec20260324>
- [22] U.S. Department of Labor, Employment and Training Administration. (2026, February 13). *AI Literacy Framework* (Training and Employment Notice No. 07-25). <https://www.dol.gov/sites/dolgov/files/ETA/advisories/TEN/2025/TEN%2007-25/>
- [23] Hinton, G. (2025, December). Interview with CNN. Reported in: Fortune. (2025, December 28). Geoffrey Hinton: AI will be even better in 2026 and replace human workers. <https://fortune.com/2025/12/28/geoffrey-hinton-godfather-of-ai-2026-prediction-human-worker-replacement/>
- [24] Revesz, R. (2024, December 27). Godfather of AI shortens odds of technology wiping out humanity over next 30 years. *The Guardian*. <https://www.theguardian.com/technology/2024/dec/27/godfather-of-ai-raises-odds-of-the-technology-wiping-out-humanity-over-next-30-years>
- [25] Hinton, G. (2026, April). Address at the 2026 Digital World Conference, Geneva. Reported in multiple outlets including France 24 and UN News.
- [26] Murray, C. (2026, February 9). Anthropic AI safety researcher warns of world "in peril" in resignation. *Forbes*. <https://www.forbes.com/sites/conormurray/2026/02/09/anthropic-ai-safety-researcher-warns-of-world-in-peril-in-resignation/>
- [27] Anthropic. (2026, February 24). *Responsible Scaling Policy version 3.0*. <https://www.anthropic.com/responsible-scaling-policy/rsp-v3-0>
- [28] Amodei, D. (2026, February 26). *Statement from Dario Amodei on our discussions with the Department of War*. Anthropic. <https://www.anthropic.com/news/statement-department-of-war>
- [29] Reuters. (2026, February 27). Trump directs US agencies to toss Anthropic's AI as Pentagon calls startup a supply risk. <https://www.reuters.com/world/us/trump-says-he-is-directing-federal-agencies-cease-use-anthropic-technology-2026-02-27/>
- [30] *Anthropic PBC v. U.S. Department of War et al.*, No. 3:26-cv-01996 (N.D. Cal., filed March 9, 2026). <https://clearinghouse.net/case/47876/>

- [31] *Anthropic PBC v. U.S. Department of War et al.*, No. 3:26-cv-01996, Doc. 134 (N.D. Cal. March 26, 2026). <https://law.justia.com/cases/federal/district-courts/california/candce/3%3A2026cv01996/465515/134/>
- [32] *Anthropic PBC v. U.S. Department of War et al.*, No. 26-1049, per curiam order (D.C. Cir. April 8, 2026).
- [33] Perrigo, B. (2026, February 24). Exclusive: Anthropic drops flagship safety pledge. *TIME*. <https://time.com/7380854/exclusive-anthropic-drops-flagship-safety-pledge/>
- [34] European Parliament and Council of the European Union. (2024). Regulation (EU) 2024/1689 of the European Parliament and of the Council on artificial intelligence (Artificial Intelligence Act). *Official Journal of the European Union*.
- [35] Council of the European Union. (2026, May 7). *Artificial intelligence: Council and Parliament agree to simplify and streamline rules* [Press release]. <https://www.consilium.europa.eu/en/press/press-releases/2026/05/07/artificial-intelligence-council-and-parliament-agree-to-simplify-and-streamline-rules/>

The Other AI: Audio Briefings on Augmented Intelligence and AI Governance
#AIassisted using the HAIA Ecosystem



The Other AI: Audio Briefings on Augmented Intelligence and AI Governance

Spotify · Apple Podcasts · Amazon Music · YouTube

#AIassisted using the HAIA Ecosystem